NGINX — Real-Time API Latency Report

Why API Performance Matters

APIs lie at the very heart of modern applications and evolving digital architectures. In today's landscape, where the barrier of switching to a digital competitor is very low, it is of the upmost importance for consumers to have positive experiences. This is ultimately driven by responsive, healthy, and adaptable APIs. If you get this right, and your API call is faster than your competitor's, developers will choose you.

However, it's a major challenge for most businesses to process API calls in as near to real time as possible. According to the IDC report *APIs — The Determining Agents Between Success or Failure of Digital Business*, over 90% of organizations expect a latency of under 50 milliseconds, while almost 60% expect latency of 20 milliseconds or less. At NGINX, we've used this data, together with some end-to-end analysis of the API lifecycle, to define a real-time API as one with latency of 30ms or less. (Latency is defined as the amount of time it takes for your API infrastructure to respond to an API call – from the moment a request arrives at the API gateway to when the first byte of a response is returned to the client.)

So, how do your APIs measure up? Are they already fast enough to be considered real time, or do they need to improve? Does your product feel a bit sluggish, but you can't quite place why that is? Maybe you don't know for sure what your API latency looks like? Whether you're using an API as the interface for microservices deployments, building a revenue stream with an external API, or something totally new, we're here to help.

Your API Performance

We have run a simple HTTP benchmark using the query parameters you specified on each of the target API endpoints you listed and created an Hdr Histogram graph that shows the latency of your API endpoints. Ideally, the latency at the 99th percentile (**99%** on the graph) is less than 30ms for your API to be considered real time.

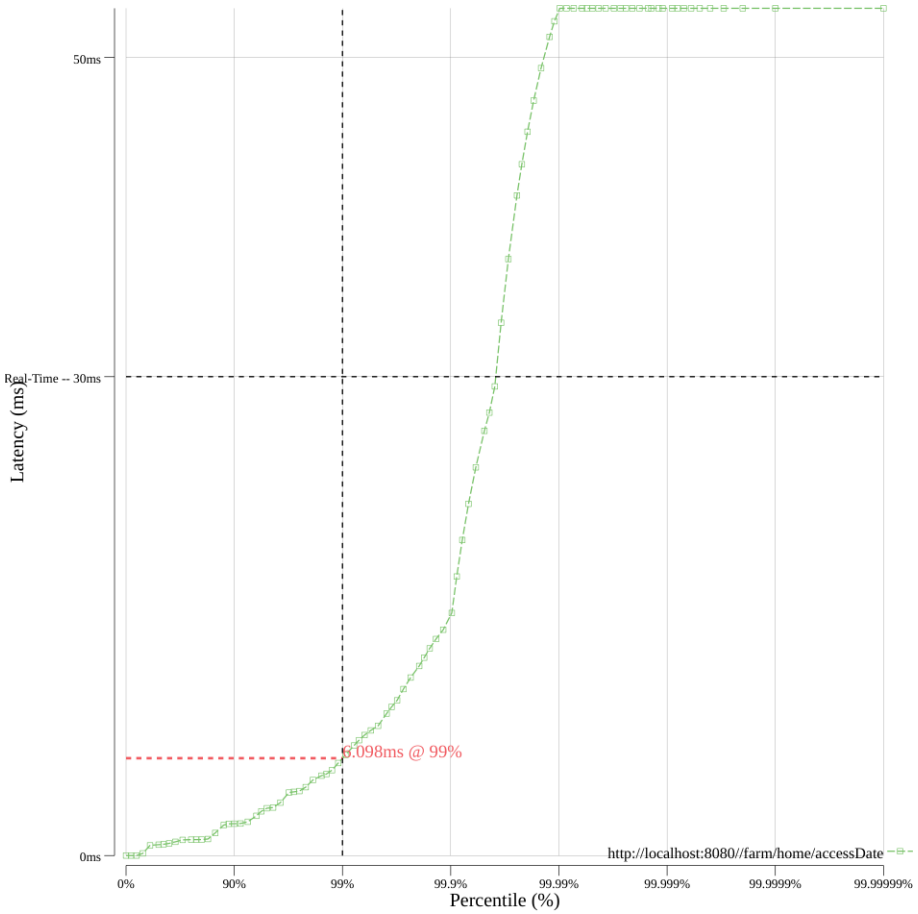

Is your API's latency below 30ms? We can help you improve it no matter where it is!

Learn more, talk to an NGINX expert, and discover how NGINX can help you on your journey towards real-time APIs at https://www.nginx.com/real-time-api